

1 Overview

Below I outline the various tests I ran in the competition, each test has it's own approach and thus its own overview, data, learning, and model selection components.

1 No Learning

1 Overview

The first submissions I made included a random output selection, which returned 49% correctness (surprise surprise), and an all zeros submission, which garnered an 83% accuracy. This test gave me information about the distribution of recession to no-recession reports. Thus off the bat, and assuming that the hidden test set is of the same distribution (we were told that it is), I should expect for a good data train, that the learning algorithm classifies 83% of the data into the no-recession category.

2 Data Manipulation

3 Learning Algorithm

4 Model Selection

SVM implementation using polynomials of varying degrees:

```
data read
fitting data
print(len(clf.dual_coef_))
1
print(len(clf.support_vectors_))
728
predicting results
train_error 0.0
classified 0.811613295501 % as non-recession for C= 2
fitting data
print(len(clf.dual_coef_))
1
print(len(clf.support_vectors_))
688
predicting results
train_error 0.0
classified 0.806039724362 % as non-recession for C= 3
fitting data
print(len(clf.dual_coef_))
1
print(len(clf.support_vectors_))
700
predicting results
train_error 0.00025
classified 0.799047426024 % as non-recession for C= 5
```

DEGREE


Percent classified as No-Recession (should be close to 83%)

```

39     for C in [2,3,5]:
40         X1 = x_train
41         results_to_num = y_train
42         .
43
44         clf = svm.SVC(
45             kernel='poly',
46             C=1,
47             degree=C,
48             shrinking=False,
49             gamma=1,
50             coef0=1)
51
52         print 'fitting data'
53         clf.fit(X1, results_to_num)
54

```

Submitted deg 2, received 89%:

#	Δ12h	Team Name	Score ?	Entries	Last Submission UTC (Best - Las
1	new	fboemer	0.93450	1	Fri, 06 Feb 2015 02:20:28
2	new	crf_wtf	0.89475	3	Fri, 06 Feb 2015 11:07:11
<p>Your Best Entry ↑</p> <p>Top Ten!</p> <p>You made the top ten by improving your score by 0.05875.</p> <p>You just moved up 1 position on the leaderboard. Tweet this!</p>					
3	new	Eman 	0.83600	5	Fri, 06 Feb 2015 00:53:22 (-0

[Download raw data](#)

SVM implementation using radial basis functions of varying degrees (2,3,5 deg):

```

[0 0 0 ..., 1 0 0]
data read
fitting data
print(len(clf.dual_coef_))
1
print(len(clf.support_vectors_))
3991
predicting results
train_error 0.0005
classified 0.999898662343 % as non-recession for C= 2
fitting data
print(len(clf.dual_coef_))
1
print(len(clf.support_vectors_))
3991
predicting results
train_error 0.0005
classified 0.999898662343 % as non-recession for C= 3
fitting data
print(len(clf.dual_coef_))
1
print(len(clf.support_vectors_))
3991
predicting results
train_error 0.0005
classified 0.999898662343 % as non-recession for C= 5

```

```

38
39     for C in [2,3,5]:
40         X1 = x_train
41         results_to_num = y_train
42
43         clf = svm.SVC(C=1.0,
44                       cache_size=200,
45                       class_weight=None,
46                       coef0=0.0,
47                       degree=C,
48                       gamma=0.0,
49                       kernel='rbf',
50                       max_iter=-1,
51                       probability=False,
52                       random_state=None,
53                       shrinking=True,
54                       tol=0.001,
55                       verbose=False)
56
57

```

This had poor results, leaving most of the points classified as non-recession.


Next tried was Linear Support Vector Classification of varying tolerances (.0001,.001,.01) and it had the various probabilities:

Decided to submit both 0.001 and 0.001:

.0001:

2
new
crf_wtf
0.89775
4
Fri, 06 Feb 2015 11:21:55

Your Best Entry ↑
Top Ten!
You made the top ten by improving your score by 0.00300.

 Tweet this!

0.001 only got 86% out of sample error.

```
[0 0 0 ..., 1 0 0]
data read
fitting data
predicting results
train_error 0.0625
classified 0.872517227402 % as non-recession for C= 0.0001
fitting data
predicting results
train_error 0.083
classified 0.759424402108 % as non-recession for C= 0.001
fitting data
predicting results
train_error 0.06675
classified 0.775030401297 % as non-recession for C= 0.01
```

Try: use lasso to make the weights sparse, then use SVM rbf with kernel trick, lasso has convergence issues on this run:

```
fitting data
E:\Anaconda\lib\site-packages\sklearn\linear_model\coordinate_descent.py:490: ConvergenceWarning:
ConvergenceWarning)
  warnings.warn('ConvergenceWarning: %s' % msg, ConvergenceWarning)
48     clf = linear_model.Lasso(alpha=C, copy_X=True, fit_intercept=True, max_iter=1000,
49         normalize=False, positive=False, precompute='auto', tol=0.0001,
50         warm_start=False)
```

Tried decision tree and got a very promising result:

```
C:\Users\dfoor\Documents\caltech\cs155\kaggle>kaggle.py
k3.csv
[0 0 0 ..., 1 0 0]
data read
fitting data
predicting results
train_error 0.0
classified 0.832590190515 % as non-recession for C= 0.0001
C:\Users\dfoor\Documents\caltech\cs155\kaggle>gvim output_decision_tree.csv
```

it classified 83.2 % as belonging to the non-recession, which, is the same distribution as the test set.

Tried random forest using cross validation and selected the model with a minimum group size of 5:

```
clf = RandomForestClassifier(n_estimators=30, max_depth=None,
    min_samples_split=C, random_state=0)
scores = cross_val_score(clf, X, y)
print scores.mean()
```

```

C:\Users\dfoor\Documents\caltech\cs155\kaggle>kaggle.py
using pruned
[0 0 0 ..., 1 0 0]
data read
0.861250305778
0.872250811531
fitting data 1
predicting results
train_error 0.00125
full classified 0.927239562221 % as non-recession for C= 3
0.856497051774
0.873247310279
fitting data 1
predicting results
train_error 0.00225
full classified 0.92521280908 % as non-recession for C= 4
0.857998553276
0.875000562782
fitting data 1
predicting results
train_error 0.00275
full classified 0.922679367653 % as non-recession for C= 5
0.855248426838
0.873750437094
fitting data 1
predicting results
train_error 0.00275
full classified 0.918321848399 % as non-recession for C= 6
0.856749928339
0.86924743334
fitting data 1
predicting results
train_error 0.0045
full classified 0.920956627483 % as non-recession for C= 7

```

4	↓1	crf_wtf	0.89875	15	Sun, 08 Feb 2015 09:31:04 (-0.0)
---	----	---------	---------	----	----------------------------------

Your Best Entry ↑

Your submission scored **0.88825**, which is not an improvement of your best score. Keep trying

Woo hoo. 90%

Got smart and started to use crossvalidation. I tried multiple SVM implementations (radial basis functions, stochastic gradient descent on linear model, polynomial models), and found that tuning C to a relaxed .0001 lead to the best performance.

```

7 #run svm over the sparse set
8 for C in [.00001,.0001,.0010]:
9     []
10     clf = svm.LinearSVC(
11         fit_intercept=True,
12         C=C)
13
14     this_scores = cross_validation.cross_val_score(clf, X, y, n_jobs=-1)
15     print 'C = ', C, 'score ', np.mean(this_scores)
16 #sys.exit()
17

```

```
C:\Users\dfloor\Documents\caltech\cs155\kaggle>kaggle.py
using pruned
[0 0 0 ..., 1 0 0]
data read
C = 1e-05 score 0.888243565905
C = 0.0001 score 0.90249244747
C = 0.001 score 0.892744693719
fitting data 1
predicting results
train_error 0.044
full classified 0.84627077422 % as non-recession for C= 10
```

5	↓1	crf_wtf	0.91200	20	Mon, 09 Feb 2015 08:57
---	----	---------	---------	----	------------------------

Your Best Entry ↑

Top Ten!

You made the top ten by improving your score by 0.01325.

You just moved up 1 position on the leaderboard. [Tweet this!](#)

Tried random forest classifier in CV:

Best was around 88%.

```
data read
C = 1 score 0.872001311656
C = 2 score 0.877245686466
C = 3 score 0.87550143847
C = 4 score 0.878248563406
C = 5 score 0.875498436968
C = 6 score 0.87249280765
C = 7 score 0.878748688719
fitting data 1
```

Tried AdaBoostClassifier:

Best I could get was around 89%:

```
data read
C = 2 score 0.897994696346
```

AdaBoostRegressor:

```
data read
C = linear score 0.273730008733
C = square score -0.426663412456
C = exponential score -0.311604886007
```

Tried the various linear_model classes. Each that didn't error out are shown with their first-pass CV accuracy. It looks like RidgeClassifier would be a good place to look:

```

1 -0.0011527622444 C = 1 LassoLars score
2 -1.17623114973 C = 1 Lars score
3 -1.46167736809 C = 1 PassiveAggressiveRegressor score
4 0.0277642308543 C = 1 ARDRegression score
5 0.136988060211 C = 1 Lasso score
6 0.162882490629 C = 1 ElasticNet score
7 0.162882490629 C = 1 ElasticNet score
8 0.241943439458 C = 1 OrthogonalMatchingPursuit score
9 0.242919934877 C = 1 LarsCV score
10 0.245637533578 C = 1 OrthogonalMatchingPursuitCV score
11 0.261153894626 C = 1 LinearRegression score
12 0.261171212242 C = 1 Ridge score
13 0.261326752113 C = 1 RidgeCV score
14 0.325136732245 C = 1 LassoLarsCV score
15 0.328784747943 C = 1 LassoLarsIC score
16 0.339102000196 C = 1 BayesianRidge score
17 0.339102000196 C = 1 BayesianRidge score
18 0.784024279152 C = 1 Perceptron score
19 0.842749421085 C = 1 PassiveAggressiveClassifier score
20 0.853739421581 C = 1 SGDClassifier score
21 0.853739421581 C = 1 SGDClassifier score
22 0.882741812277 C = 1 LogisticRegression score
23 0.882742562653 C = 1 RidgeClassifier score
24 0.882742562653 C = 1 RidgeClassifierCV score

```

Couldn't get Ridge any better.

Tried neural nets;

```

186 [] clf=neural_network.BernoulliRBM(n_components=1000, learning_rate=0.1, batch_size=10, n_iter=
fitting data 1
predicting results
train_error 0.23575
full classified 0.0769152817187 % as non-recession for C= 10

```

Used grid search to explore the various parameter spaces. Looks like rbf gets possibly better than the current victor—the linear $C=.0001$ model.

```

Grid scores on development set:
()
0.000 (+/-0.000) for {'kernel': 'rbf', 'C': 0.0001, 'gamma': 0.001}
0.000 (+/-0.000) for {'kernel': 'rbf', 'C': 0.0001, 'gamma': 0.0001}
0.000 (+/-0.000) for {'kernel': 'rbf', 'C': 1, 'gamma': 0.001}
0.048 (+/-0.011) for {'kernel': 'rbf', 'C': 1, 'gamma': 0.0001}
0.000 (+/-0.000) for {'kernel': 'rbf', 'C': 10, 'gamma': 0.001}
0.086 (+/-0.019) for {'kernel': 'rbf', 'C': 10, 'gamma': 0.0001}
0.000 (+/-0.000) for {'kernel': 'rbf', 'C': 100, 'gamma': 0.001}
0.086 (+/-0.019) for {'kernel': 'rbf', 'C': 100, 'gamma': 0.0001}
0.000 (+/-0.000) for {'kernel': 'rbf', 'C': 1000, 'gamma': 0.001}
0.086 (+/-0.019) for {'kernel': 'rbf', 'C': 1000, 'gamma': 0.0001}
0.650 (+/-0.021) for {'kernel': 'poly', 'C': 1, 'degree': 2}
0.605 (+/-0.023) for {'kernel': 'poly', 'C': 1, 'degree': 3}
0.558 (+/-0.014) for {'kernel': 'poly', 'C': 1, 'degree': 4}
0.650 (+/-0.021) for {'kernel': 'poly', 'C': 10, 'degree': 2}
0.605 (+/-0.023) for {'kernel': 'poly', 'C': 10, 'degree': 3}
0.558 (+/-0.014) for {'kernel': 'poly', 'C': 10, 'degree': 4}
0.650 (+/-0.021) for {'kernel': 'poly', 'C': 100, 'degree': 2}
0.605 (+/-0.023) for {'kernel': 'poly', 'C': 100, 'degree': 3}
0.558 (+/-0.014) for {'kernel': 'poly', 'C': 100, 'degree': 4}
0.650 (+/-0.021) for {'kernel': 'poly', 'C': 1000, 'degree': 2}
0.605 (+/-0.023) for {'kernel': 'poly', 'C': 1000, 'degree': 3}
0.558 (+/-0.014) for {'kernel': 'poly', 'C': 1000, 'degree': 4}
0.546 (+/-0.020) for {'kernel': 'linear', 'C': 0.0001}
0.635 (+/-0.014) for {'kernel': 'linear', 'C': 0.01}
0.629 (+/-0.020) for {'kernel': 'linear', 'C': 1}
0.629 (+/-0.020) for {'kernel': 'linear', 'C': 10}
0.629 (+/-0.020) for {'kernel': 'linear', 'C': 100}
0.629 (+/-0.020) for {'kernel': 'linear', 'C': 1000}
()
Detailed classification report:
()
The model is trained on the full development set.
The scores are computed on the full evaluation set.
()

```

	precision	recall	f1-score	support
0.0	0.92	0.92	0.92	1630
1.0	0.64	0.64	0.64	370
avg / total	0.82	0.82	0.82	2000

clearly, nope—and it matches the CV score:

6	—	crf_wtf	0.91200	21	Tue, 10 Feb 2015 07:4
Your Best Entry ↑					
Your submission scored 0.83600 , which is not an improvement of your best score. Keep trying					

heh heh.

I took the results from the test set and included them in my training set. Because my regularizer is pretty loose, and because I scored better than 1/2 on the test set, I was able to increase the score—I think. I will try this recursively until something theoretical breaks it.

5	↑1	crf_wtf	0.91500	22	Tue, 10 Feb 2015 08:0
---	----	---------	---------	----	-----------------------

Your Best Entry ↑

Top Ten!

You made the top ten by improving your score by 0.00300.

You just moved up 1 position on the leaderboard. [Tweet this!](#)

Second time I tried this, it broke:

6	↓1	crf_wtf	0.91500	23	Tue, 10 Feb 2015 09:00
---	----	---------	---------	----	------------------------

Your Best Entry ↑

Your submission scored **0.91325**, which is not an improvement of your best score. Keep trying

Okay—using adaboost, I got a better result and still had training error, so I'll try more classifiers:

8	↓2	crf_wtf	0.91775	29	Wed, 11 Feb 2015 08:0
---	----	---------	---------	----	-----------------------

Your Best Entry ↑

Top Ten!

You made the top ten by improving your score by 0.00275.

You just moved up 1 position on the leaderboard. [Tweet this!](#)

GradientBoostClassifier returned even better error at 93% and with a tuned max depth of 3:

```

126
127     clf = GradientBoostingClassifier(n_estimators=1000, learning_rate=1.0,
128                                     max_depth=3, random_state=0)
129

```

6	—	crf_wtf	0.93125	30	Wed, 11 Feb 2015 08:26
---	---	---------	---------	----	------------------------

5 Conclusion